EXAM ASSOCIATE CLOUD ENGINEER TOPIC 1 QUESTION 81 DISCUSSION

Actual exam question from Google's Associate Cloud Engineer
Question #: 81
Topic #: 1
[All Associate Cloud Engineer Questions]

You are operating a Google Kubernetes Engine (GKE) cluster for your company where different teams can run non-production workloads. Your Machine Learning
(ML) team needs access to Nvidia Tesla P100 GPUs to train their models. You want to minimize effort and cost. What should you do?

A. Ask your ML team to add the ג€accelerator: gpuג€ annotation to their pod specification.

B. Recreate all the nodes of the GKE cluster to enable GPUs on all of them.

C. Create your own Kubernetes cluster on top of Compute Engine with nodes that have GPUs. Dedicate this cluster to your ML team.

D. Add a new, GPU-enabled, node pool to the GKE cluster. Ask your ML team to add the cloud.google.com/gke -accelerator: nvidia-tesla-p100 nodeSelector to their pod specification.

Show Suggested Answer

by John_Iam at *June 3, 2020, 9:22 p.m.*

Comments

Type your comment...

Submit

John_Iam **Highly Voted** 4 years, 4 months ago

D is the correct answer.

https://cloud.google.com/kubernetes-engine/docs/how-to/gpus

upvoted 47 times

tablet444 4 years, 2 months ago

the documentation states "Limitations
Before using GPUs on GKE, keep in mind the following limitations:

You cannot add GPUs to existing node pools.
GPU nodes cannot be live migrated during maintenance events."

upvoted 10 times

nightflyer 3 years, 10 months ago

In this case it is about adding a GPU enabled node pool not a GPU to an existing node-pool

upvoted 16 times

fragment137 1 year, 10 months ago

You're correct that D says that, except that the question also says to use the most cost-effective method. Two node-pools would be more expensive than rebuilding the current one with GPU enabled.

upvoted 1 times

Gulithor 1 year, 9 months ago

It also says to minimize effort, wouldn't recreating all the pools take way longer than just adding 1?

upvoted 1 times

glam **Highly Voted** 4 years ago

D. Add a new, GPU-enabled, node pool to the GKE cluster. Ask your ML team to add the cloud.google.com/gke -accelerator: nvidia-tesla-p100 nodeSelector to their pod specification.

upvoted 15 times

BAofBK **Most Recent** 11 months, 2 weeks ago

The correct answer is D

upvoted 1 times

trainingexam 1 year, 3 months ago

Selected Answer: D

D is the correct answer.

upvoted 1 times

sabrinakloud 1 year, 6 months ago

Selected Answer: D

Option D is good

upvoted 1 times

sakdip66 1 year, 6 months ago

Selected Answer: D

Creating new node pool w/ GPU-enabled instances is cost - saving solution. This way ML team workload will GPU instance for their ML and other team workload will run smoothly

upvoted 2 times

Prat25200607 1 year, 6 months ago

Selected Answer: D

D makes more cost effective

upvoted 1 times

Buruguduystunstugudunstuy 1 year, 8 months ago

Selected Answer: D

Answer D. Add a new, GPU-enabled, node pool to the GKE cluster. Ask your ML team to add the cloud.google.com/gke-accelerator: nvidia-tesla-p100 nodeSelector to their pod specification.

Adding a new node pool with GPUs is the best option because it allows for a separate set of nodes that can be specifically allocated to workloads that require GPU acceleration, such as the Machine Learning (ML) team's workloads. This approach will not affect other workloads running on the original nodes, keeping the costs low and the overall cluster performance stable

stable.

upvoted 6 times

cslince 1 year, 10 months ago

Selected Answer: D

D is the correct answer.

upvoted 1 times

leogor 1 year, 11 months ago

D is correct

upvoted 1 times

raghu09 2 years, 1 month ago

Selected Answer: D

D is correct

upvoted 1 times

iadarsh 2 years, 1 month ago

Selected Answer: D

D is correct
Because if you create entirely new node pool then its not cost efficient and also the pods which not require that much high GPU is get scheduled into it. So instead of that add a new node pool with GPU and in the pod YAML file mention the node affinity to get scheduled into the GPU enabled node pool.

upvoted 2 times

AzureDP900 2 years, 4 months ago

By looking at all answers first 3 can be eliminated without any second thought. D is correct.

upvoted 2 times

haroldbenites 2 years, 4 months ago

Go for D

upvoted 1 times

LaxmanTiwari 2 years, 5 months ago

Selected Answer: D

In this case it is about adding a GPU enabled node pool not a GPU to an existing node-pool

upvoted 2 times

luciorifa 2 years, 8 months ago

Selected Answer: D

D is the correct answer

upvoted 1 times

jaffarali 2 years, 10 months ago

Selected Answer: D

D would be the right option when there is possibility to add GPUs without recreating the nodes.

upvoted 3 times

Load full discussion...

Social Media

Facebook , Twitter
YouTube , Reddit
Pinterest

We are the biggest and most updated IT certification exam material website.
Using our own resources, we strive to strengthen the IT professionals community for free.