



- Expert Verified, Online, Free.

MENU



Google Discussions



Exam Associate Cloud Engineer All Questions

View all questions & answers for the Associate Cloud Engineer exam

Go to Exam

EXAM ASSOCIATE CLOUD ENGINEER TOPIC 1 QUESTION 139 DISCUSSION

Actual exam question from Google's Associate Cloud Engineer

Question #: 139

Topic #: 1

[\[All Associate Cloud Engineer Questions\]](#)

The core business of your company is to rent out construction equipment at large scale. All the equipment that is being rented out has been equipped with multiple sensors that send event information every few seconds. These signals can vary from engine status, distance traveled, fuel level, and more. Customers are billed based on the consumption monitored by these sensors. You expect high throughput `` up to thousands of events per hour per device `` and need to retrieve consistent data based on the time of the event. Storing and retrieving individual signals should be atomic. What should you do?




- A. Create a file in Cloud Storage per device and append new data to that file.
- B. Create a file in Cloud Filestore per device and append new data to that file.
- C. Ingest the data into Datastore. Store data in an entity group based on the device.
- D. Ingest the data into Cloud Bigtable. Create a row key based on the event timestamp.

Show Suggested Answer

by [filco72](#) at Aug. 11, 2020, 2:14 p.m.

Comments

Type your comment...

  **hiteshrup** Highly Voted  4 years, 1 month ago

Answer: D

Keyword need to look for

- "High Throughput",
- "Consistent",
- "Property based data insert/fetch like engine status, distance traveled, fuel level, and more." which can be designed in column,
- "Large Scale Customer Base + Each Customer has multiple sensor which send event in seconds" This will go for per byte situation,
- Export data based on the time of the event.
- Atomic

o BigTable will fit all requirement.

o DataStore is not fully Atomic

o CloudStorage is not a option where we can export data based on time of event. We need another solution to do that

o Firestore can be used with MobileSDK.

So go with Option D: Big Table

   upvoted 48 times

  **[Removed]** 4 years, 1 month ago

Its File store not firestore. But still, the argument is correct here as file store will not be automatic like cloud store that no SQL.



If it would be Firestore I would have gone with it, then big data for throughout v/s cost.

   upvoted 4 times

  **Hjameel** Highly Voted  4 years, 2 months ago

D is the best answer , Cloud Bigtable

   upvoted 8 times

  **har_riy** 3 years, 8 months ago

Simple analogy.

Information every few seconds --> Time Series --> Big Table

   upvoted 6 times

  **SHAAHIBHUSHANAWS** Most Recent  10 months ago

Selected Answer: D



<https://cloud.google.com/bigtable/docs/overview>

   upvoted 1 times

  **Linhtinh603** 10 months, 3 weeks ago

Answer D is the best for high throughput and IoT, but I concern about creating a row key based on the event timestamp, will it be leading a hotspots issue?

   upvoted 2 times

  **Captain1212** 1 year, 1 month ago

Selected Answer: D

D is the right answer, as its can help with Automatically

   upvoted 1 times

  **Tossha** 1 year, 8 months ago

ReadModifyWriteRow requests are atomic:

<https://cloud.google.com/bigtable/docs/writes>

   upvoted 1 times

  **xaqanik** 1 year, 8 months ago

it say any type of information. it means there can be an image file for instance. so, Bigtable is the best fit for this scenario

   upvoted 1 times

  **alex000** 1 year, 9 months ago

Selected Answer: C






Answer: C






key work: atomic transaction






<https://cloud.google.com/datastore/docs/concepts/overview>






   upvoted 2 times






   upvoted 2 times






  **[Removed]** 1 year, 6 months ago
high throughput == Bigtable, hence the answer is D
   upvoted 1 times






  **Cornholio_LMC** 2 years ago
had this question today
   upvoted 5 times






  **tomis2** 2 years, 3 months ago
Selected Answer: D
Timeseries + IoT = Bigtable
   upvoted 4 times






  **AzureDP900** 2 years, 4 months ago
This is related to IoT , it is no sql means BigTable. D is right
   upvoted 2 times






  **pfabio** 2 years, 5 months ago
Selected Answer: D
D - Ideal for use cases such as personalization, ad tech, fintech, digital media, and IoT
<https://cloud.google.com/bigtable>
   upvoted 3 times

  **JelloMan** 2 years, 5 months ago
Selected Answer: D
Big table for the exact reasons as @hitshrurp mentioned. Time related data flowing in and out will point you directly to BigTable without reading anything else
   upvoted 1 times

  **wael_tn** 2 years, 6 months ago
Selected Answer: D
When we want to store data based on imestamp => Cloud Bigtable
   upvoted 1 times

  **Akash7** 2 years, 6 months ago
Answer is D, large streaming data = bigtable
   upvoted 1 times

  **zaxma** 2 years, 6 months ago
Selected Answer: C
I vote c with DataStore (maybe even better with firestore includes pubsub ability)
Looks like the internet agrees with me.
   upvoted 2 times

  **alaahakim** 2 years, 11 months ago
Answer: D
   upvoted 2 times

[Load full discussion...](#)



Social Media

[Facebook](#) , [Twitter](#)

[YouTube](#) , [Reddit](#)

[Pinterest](#)



We are the biggest and most updated IT certification exam material website.

Using our own resources, we strive to strengthen the IT professionals community for free.



© 2024 ExamTopics

ExamTopics doesn't offer Real Microsoft Exam Questions. ExamTopics doesn't offer Real Amazon Exam Questions. ExamTopics Materials do not contain actual questions and answers from Cisco's Certification Exams.

CFA Institute does not endorse, promote or warrant the accuracy or quality of ExamTopics. CFA® and Chartered Financial Analyst® are registered trademarks owned by CFA Institute.