



- Expert Verified, Online, Free.

MENU



Google Discussions



Exam Associate Cloud Engineer All Questions

View all questions & answers for the Associate Cloud Engineer exam

Go to Exam

EXAM ASSOCIATE CLOUD ENGINEER TOPIC 1 QUESTION 23 DISCUSSION

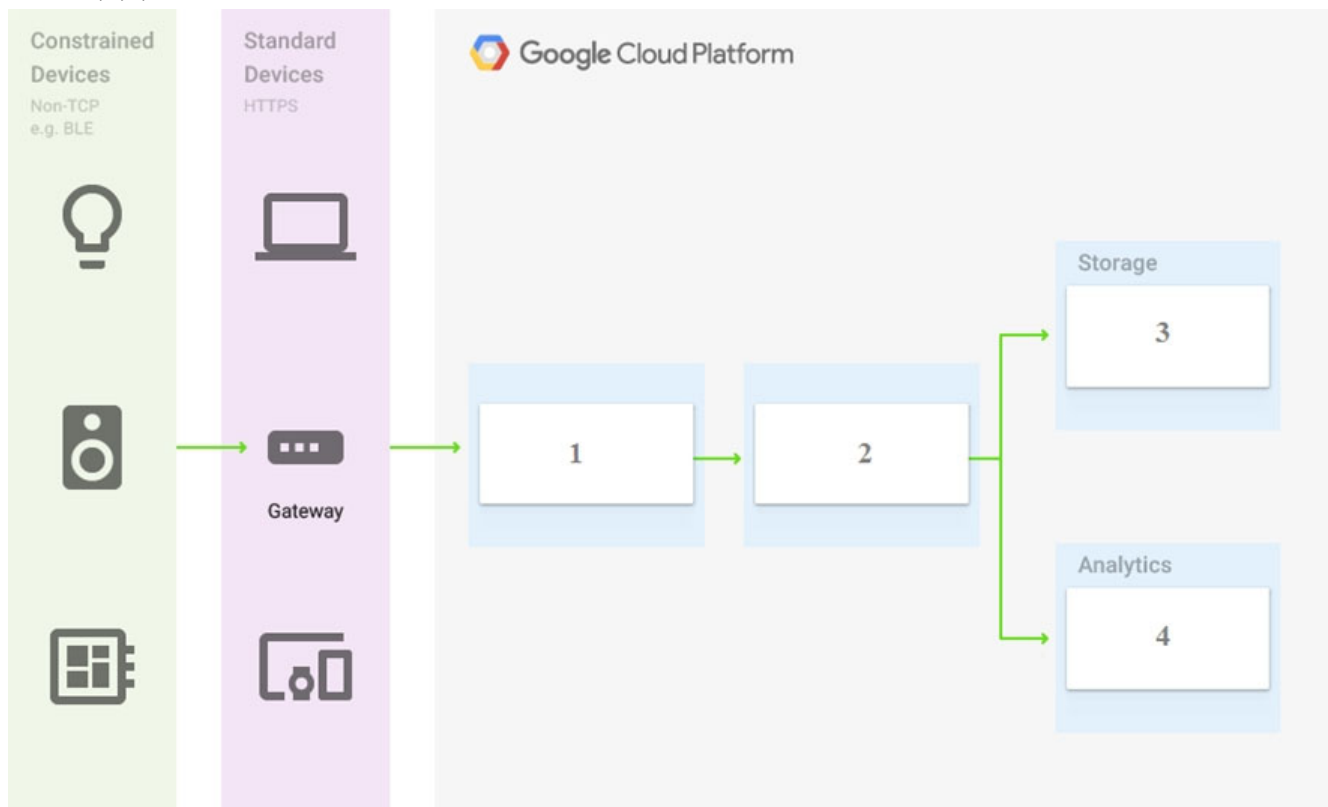
Actual exam question from Google's Associate Cloud Engineer

Question #: 23

Topic #: 1

[\[All Associate Cloud Engineer Questions\]](#)

You are building a pipeline to process time-series data. Which Google Cloud Platform services should you put in boxes 1,2,3, and 4?



- A. Cloud Pub/Sub, Cloud Dataflow, Cloud Datastore, BigQuery
- B. Firebase Messages, Cloud Pub/Sub, Cloud Spanner, BigQuery
- C. Cloud Pub/Sub, Cloud Storage, BigQuery, Cloud Bigtable
- D. Cloud Pub/Sub, Cloud Dataflow, Cloud Bigtable, BigQuery

Show Suggested Answer

by [cloudenth01](#) at June 25, 2020, 10:12 a.m.

Comments

Type your comment...

Submit

[cloudenth01](#) Highly Voted 4 years, 3 months ago

Without a doubt D.
Whenever we want to process timeseries data look for BigTable.
Also you want to perform analytics in Box 4 ..look for BigQuery

Only D provides this option.

[?](#) [?](#) [?](#) upvoted 68 times

[vloodia](#) 4 years, 3 months ago

Speaker also looks like an IoT device so D not A

[?](#) [?](#) [?](#) upvoted 3 times

[adedj99](#) 3 years, 10 months ago

are we considering bigtable as storage in here , since they expecting some storage

[?](#) [?](#) [?](#) upvoted 2 times

[Buruguduystunstugudunstuy](#) Highly Voted 1 year, 8 months ago

Selected Answer: D

The correct process for building a pipeline to process time-series data. Here's how each of the components is used:

1. Cloud Pub/Sub: receives and distributes time-series data from different sources.
2. Cloud Dataflow: processes the data by applying transformations and analytics.
3. Cloud Bigtable: stores and manages the processed data as a NoSQL database.
4. BigQuery: provides a SQL-like interface to analyze the data and extract insights.

By combining these components, you can create a scalable and reliable pipeline to process and analyze time-series data in real time.

[?](#) [?](#) [?](#) upvoted 42 times

[kenrichy](#) 1 year, 7 months ago

Many thanks Buruguduy.. for the extensive explanation you always give to your choice of answers. It's really helpful to understand the concept.

[?](#) [?](#) [?](#) upvoted 5 times

[BAofBK](#) Most Recent 11 months, 2 weeks ago

Correct answer D

[?](#) [?](#) [?](#) upvoted 1 times

[YourCloudGuru](#) 1 year ago

Selected Answer: D

The correct answer is D

This diagram shows a typical pipeline for processing time-series data:

1. ****Cloud Pub/Sub:**** A messaging service that allows you to send and receive messages between independent applications

applications.

2. **Cloud Dataflow:** A fully-managed service for transforming and processing data streams.
3. **Cloud Bigtable:** A wide-column, distributed NoSQL database that is optimized for storing and analyzing large amounts of data.
4. **BigQuery:** A fully-managed, petabyte-scale analytics data warehouse that enables businesses to analyze all their data very quickly.

In this pipeline, the following happens:

1. Time-series data is sent to Cloud Pub/Sub.
2. Cloud Dataflow reads the data from Cloud Pub/Sub and performs any necessary transformations or processing.
3. Cloud Dataflow writes the transformed data to Cloud Bigtable.
4. BigQuery queries the data in Cloud Bigtable to generate insights.

? ? ? upvoted 3 times

? ? **Captain1212** 1 year, 1 month ago

big table for 3, so d is the correct answer

? ? ? upvoted 1 times

? ? **Neha_Pallavi** 1 year, 3 months ago

D. Cloud Pub/Sub, Cloud Dataflow, Cloud Bigtable, BigQuery

? ? ? upvoted 1 times

? ? **Partha117** 1 year, 7 months ago

Selected Answer: D

Cloud Bigtable for time series data storage and BigQuery for analysis

? ? ? upvoted 1 times

? ? **csllince** 1 year, 10 months ago

Selected Answer: D

Option D

? ? ? upvoted 1 times

? ? **dennydream** 1 year, 11 months ago

I know the answer is D, but it's misleading. "Storage" would suggest cloud storage, not BigTable.

? ? ? upvoted 2 times

? ? **akshaydoifode88** 1 year, 11 months ago

You can't store timeseries data in cloud storage. Even if you store it how you are gonna access it?

? ? ? upvoted 1 times

? ? **ale_brd_111** 2 years, 1 month ago

Selected Answer: D

whenever IOT is present go with DataStore

? ? ? upvoted 1 times

? ? **RanjithK** 2 years, 3 months ago

Answer is D

? ? ? upvoted 1 times

? ? **AzureDP900** 2 years, 4 months ago

D is right

? ? ? upvoted 1 times

? ? **Pleaseletmein** 2 years, 4 months ago

D is the right answer

? ? ? upvoted 1 times

? ? **bamosk** 2 years, 4 months ago

Selected Answer: D

Look the picture posted on this page here

<https://cloud.google.com/dataflow#section-7>

? ? ? upvoted 2 times

? ? **bamosk** 2 years, 4 months ago

IoT = Unstructured data -> eliminated Datastore + Timeseries requirement = BigTable

Ingestion point = Pub/Sub (Firebase messaging is a push notification service to client)

Process data + realtime or batch = Data flow

Analytics = BigQuery

Hence D.

? ? ? upvoted 6 times

  **mplibunao** 2 years, 4 months ago

Selected Answer: D

Answer is D

   upvoted 1 times

  **haroldbenites** 2 years, 5 months ago

Go for D

   upvoted 1 times

[Load full discussion...](#)

Start Learning for free



Social Media

[Facebook](#) , [Twitter](#)

[YouTube](#) , [Reddit](#)

[Pinterest](#)



We are the biggest and most updated IT certification exam material website.

Using our own resources, we strive to strengthen the IT professionals community for free.



© 2024 ExamTopics

ExamTopics doesn't offer Real Microsoft Exam Questions. ExamTopics doesn't offer Real Amazon Exam Questions. ExamTopics Materials do not contain actual questions and answers from Cisco's Certification Exams.

CFA Institute does not endorse, promote or warrant the accuracy or quality of ExamTopics. CFA® and Chartered Financial Analyst® are registered trademarks owned by CFA Institute.