



- Expert Verified, Online, Free.

 MENU



Google Discussions



## Exam Associate Cloud Engineer All Questions

View all questions & answers for the Associate Cloud Engineer exam

Go to Exam

### EXAM ASSOCIATE CLOUD ENGINEER TOPIC 1 QUESTION 275 DISCUSSION

Actual exam question from Google's Associate Cloud Engineer

Question #: 275

Topic #: 1

[\[All Associate Cloud Engineer Questions\]](#)

You have several hundred microservice applications running in a Google Kubernetes Engine (GKE) cluster. Each microservice is a deployment with resource limits configured for each container in the deployment. You've observed that the resource limits for memory and CPU are not appropriately set for many of the microservices. You want to ensure that each microservice has right sized limits for memory and CPU. What should you do?

- A. Configure a Vertical Pod Autoscaler for each microservice.
- B. Modify the cluster's node pool machine type and choose a machine type with more memory and CPU.
- C. Configure a Horizontal Pod Autoscaler for each microservice.
- D. Configure GKE cluster autoscaling.

Show Suggested Answer

by  [RuchiMishra](#) at July 16, 2024, 12:12 p.m.

## Comments

Type your comment...

Submit

? ? jhumpamp 2 months, 3 weeks ago

Selected Answer: C

As mentioned "resource limits configured for each container in the deployment", so can not vertically increase resources/memory

Cluster scale out is not relevant here.

With Horizontal scaling, adding more pods ultimately makes sure each service have "right sized memory and CPU"

? ? ? upvoted 2 times

? ? flummoxed\_individual 2 months, 4 weeks ago

Selected Answer: A

Vertical is for more specific resources of the individual pods. Horizontal is for creating more copies of the instances (adding more pods).

? ? ? upvoted 2 times

? ? user636 3 months ago

Selected Answer: A

A seems better

? ? ? upvoted 2 times

? ? RuchiMishra 3 months, 1 week ago

Selected Answer: A

Here's why a Vertical Pod Autoscaler (VPA) is the most suitable solution for this scenario:

Right-Sizing Resources: VPA is designed to automatically adjust the resource requests and limits (CPU and memory) for pods based on their actual usage. This ensures that pods have enough resources to run efficiently without being over-provisioned, which can lead to wasted resources and higher costs.

Automated Optimization: VPA continuously monitors the resource usage of your pods and recommends optimal settings. You can choose to apply these recommendations automatically or manually, giving you flexibility and control over the process.

Microservice-Specific Tuning: By configuring a VPA for each microservice, you can fine-tune the resource allocation for each individual service based on its specific needs and usage patterns. This is more efficient than making blanket changes to the entire cluster or node pool.

? ? ? upvoted 4 times

Start Learning for free



## Social Media

Facebook , Twitter

YouTube , Reddit

Pinterest



We are the biggest and most updated IT certification exam material website.

Using our own resources, we strive to strengthen the IT professionals community for free.



© 2024 ExamTopics

ExamTopics doesn't offer Real Microsoft Exam Questions. ExamTopics doesn't offer Real Amazon Exam Questions. ExamTopics Materials do not contain actual questions and answers from Cisco's Certification Exams.

CFA Institute does not endorse, promote or warrant the accuracy or quality of ExamTopics. CFA® and Chartered Financial Analyst® are registered trademarks owned by CFA Institute.